

ANSIBLE

Serverless + Ansible

Ryan Scott Brown
Senior Software Engineer, Ansible Core

ANSIBLE

Serverless Advanced Cloud + Ansible

Ryan Scott Brown
Senior Software Engineer, Ansible Core

- Bias review
- Ansible Philosophy
- Serverless Overview
 - Why Serverless?
- Ansible for Cloud Resources
- Fake hosts, roles, and practices
- Deployments with Ansible
 - Serverless Framework & Ansible
 - All Ansible
 - CloudFormation & Ansible
- Coming Soon
 - AWS SAM
 - More modules!

Biases

- I work for Ansible
- Public cloud fan
- Primarily focused on Linux infra

Ansible Philosophy

Automation for the People

Simple Wins

- Simple, powerful automation
- Play well with others
- Work for users
- Listen to the community

- Modules to cover different providers
- Dynamic inventory for changing resources
- Smart credential handling
- Split abstraction
 - Control-plane “on the cloud”
 - Instance level “in the cloud”
- Bring tools you already use

Serverless


So I don't need orchestration, right?

For Developers

- Pay per use with fine-grained metering
- Autoscaling driven by demand
- Invocations create their own new environment
- Fewer barriers to new development

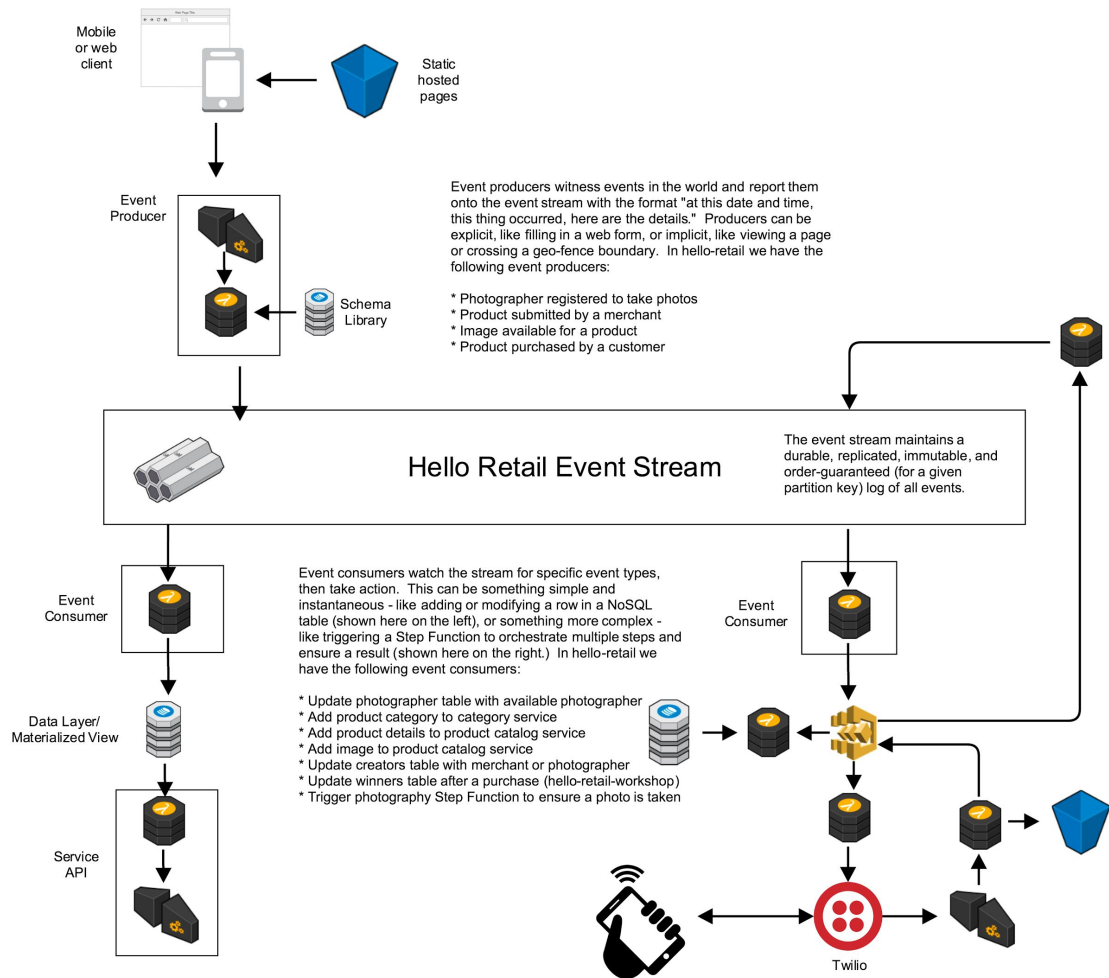
For Operators

- Autoscaling driven by demand
- Everything is behind an API
- Consistent, repeatable environment
- Dev/Prod parity



AWS Lambda, API Gateway, DynamoDB, and Friends*

- Narrowly focused tools connected via APIs
- Replace custom development with configuration
- Free up people to spend more time on other tasks

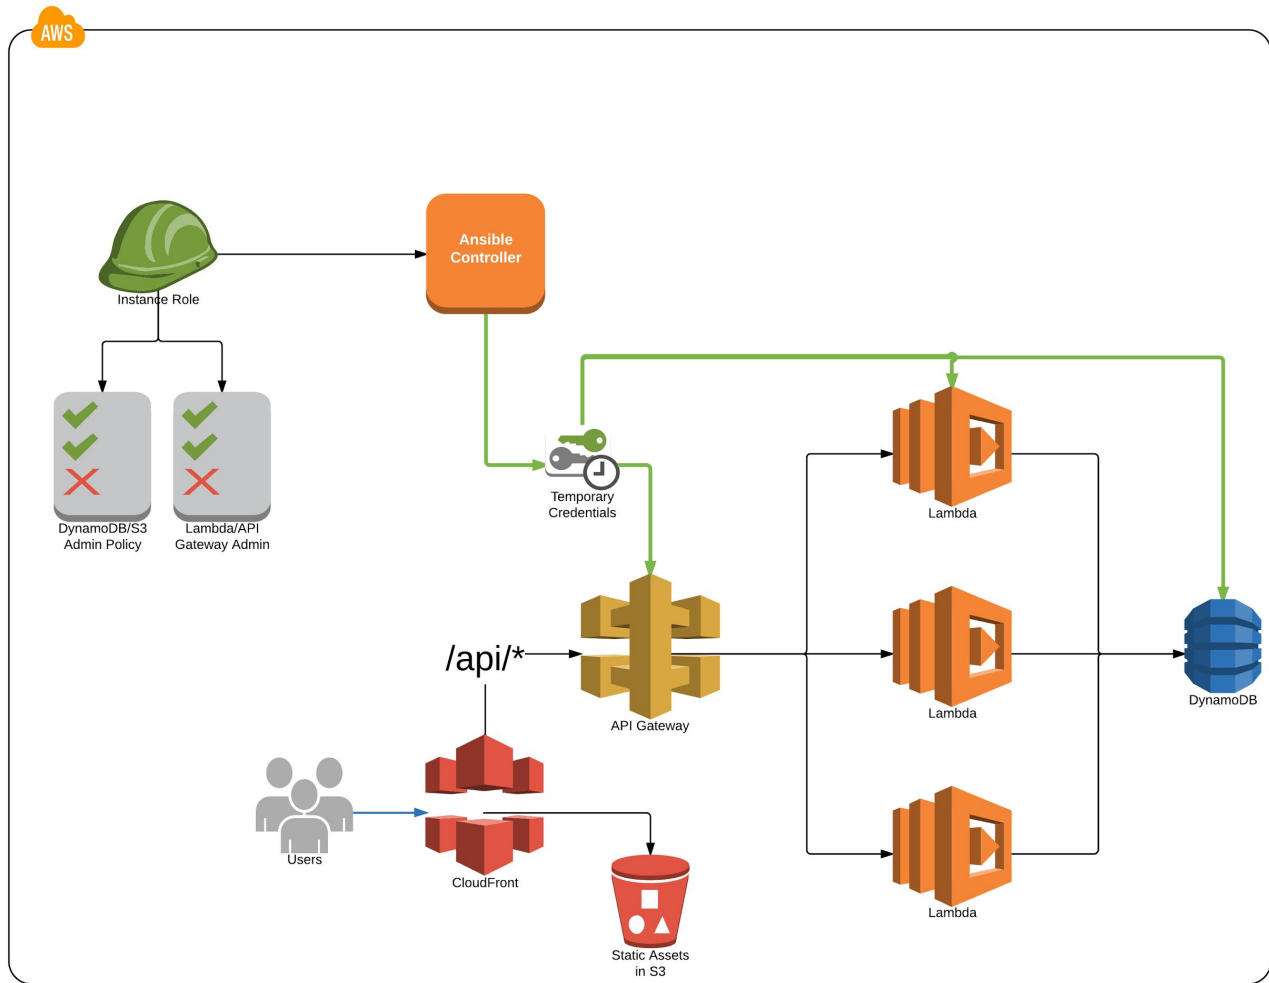


Where is Ansible?

- Ansible control can be done from anywhere with API access
 - Tower in your datacenter or EC2
 - Remote node in EC2
- Credentialed as an IAM user or instance role
- Serverless solutions trade custom components for services
 - More varied APIs
 - More configuration spread across services

AWS Lambda, API Gateway, DynamoDB, and Friends*

- Narrowly focused tools connected via **APIs**
- Replace custom development with **configuration**
- Free up people to spend more time on other tasks



Ansible Support

- aws_api_gateway
- data_pipeline
- execute_lambda
- lambda
- lambda_policy
- lambda_event
- cloudwatchevent_rule
- dynamodb_table
- dynamodb_ttl
- elasticache
- azure_rm_functionapp

But the Inventory!

Hostless?

Virtual Hosts for Service Settings

- Supports group+host variables
- Groups are reusable for host orchestration
- Minimizes vars in the global namespace

```
sls_east_prod ansible_connection=local ansible_host=127.0.0.1
sls_west_prod ansible_connection=local ansible_host=127.0.0.1
sls_west_dev ansible_connection=local ansible_host=127.0.0.1
sls_east_dev ansible_connection=local ansible_host=127.0.0.1
```

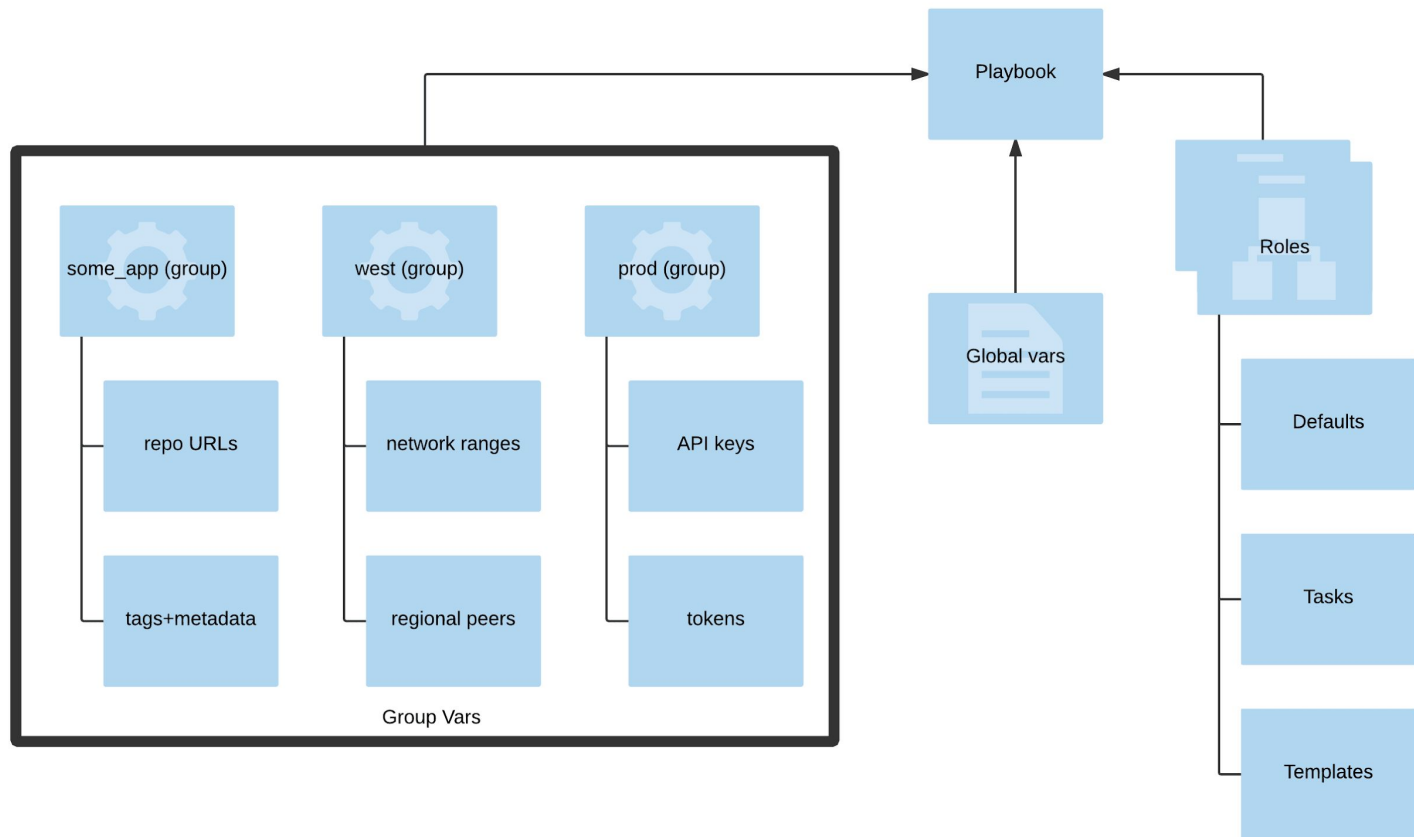
```
[prod]
```

```
sls_east_prod
sls_west_prod
```

```
[east]
```

```
sls_east_prod
sls_east_dev
```

```
region: us-east-1
azs:
- us-east-1a
- us-east-1b
- us-east-1c
- us-east-1d
- us-east-1e
peers:
  db: e1-replica-sql-{{ stage }}.{{ tld }}
tld: east.internal.myco.com.
```



Roles for Service Settings

- Use role defaults and override as needed
- Sharable between teams for common tooling
 - CloudFormation custom resources
 - Evented security/auditing functions
 - Lambda extensions to AutoScaling or ECS
- Role namespace keeps serverless-specific vars off instances

Implementation Patterns

Pattern 1

Ansible and Serverless Framework

<https://serverless.com/framework>

Best of Both Worlds

- Toolkit for (cross-provider) serverless apps
- Unified CLI and config
- Community focused exclusively on serverless tools

Why both?

- Helps to have devs with Serverless Framework experience
- Can pass control for shared resources up to Ansible

Cross-Cloud Support

- AWS Lambda
- Azure Functions
- Google Cloud Functions
- OpenWhisk
- Kubeless
- SpotInst



```
- serverless_deploy:  
  stage: prod  
  functions:  
    - createUser  
    - deleteUser  
  region: us-east-1
```

Migration Tips

- Attention to detail - who owns resources needed by both?
- Make sure data is close enough to new & old systems
- Remember Ansible can populate Serverless variable files

Exporting Variables

- name: Save out variables to YAML for sls framework
copy:
content: “{{ sls_vars | to_nice_yaml }}”
dest: “{{ project_dir }}/vars/ansible-generated.yml”

Exported Serverless Vars

```
lambda_iam_arn: arn:aws:iam:.....:role/MyExecRole  
shared_kinesis: some-event-stream
```



```
custom:
  stage: ${opt:stage, self:provider.stage, self:custom.private.stage}
  private: ${file(..secrets.yml)}
  external: ${file(vars/ansible-generated.yml)}

provider:
  name: aws
  runtime: python3.6
  environment:
    STAGE: ${self:custom.stage}

functions:
  eventWriter:
    role: ${self:custom.external.lambda_iam_arn}
    handler: eventWriter.ingest
    environment:
      STREAM_NAME:
        Fn::ImportValue: ${self:custom.exportStreamName}
    events:
      - http:
          path: event-writer
          method: post
```

Pattern 2

Ansible All-In

Fastest Start

- Ansible skills 100% transferrable
- Add to existing playbooks
- Great for mixed workloads

Build a Function

```
- lambda:
  name: sendReportMail
  zip_file: "{{ deployment_package }}"
  runtime: python2.7
  timeout: 20
  handler: handler.handler
  memory_size: 1024
  role: "{{ iam_exec_role }}"
register: lambda
```

Add HTTPS Endpoint (New in 2.4)

- `aws_api_gateway:`
 - `api_file: my-swagger-def.yml`
 - `stage: dev`

```
swagger: "2.0"
basePath: "/something"
paths:
  /:
    get:
      consumes:
        - "application/json"
      produces:
        - "application/json"
      responses:
        200:
          description: "200 response"
          schema:
            $ref: "#/definitions/Empty"
```

Cron Scheduling

- `cloudwatchevent_rule`:
 - `description`: "Send out incident summary daily"
 - `name`: `incident_summary_schedule`
 - `schedule_expression`: `'rate(1 day)'`
 - `targets`:
 - `id`: `SendSummary`
 - `arn`: `"{{ lambda.configuration.function_arn }}"`
 - `input`: `'{}'`

Pattern 3

Ansible + CloudFormation

Self-Contained Apps

- If you have CloudFormation skills already
- AWS publishes reference architectures
- Check mode now uses full ChangeSets
- Serverless Application Model?

Sample Play

- cloudformation:
 - stack_name: “{{ env }}-sls-app”
 - state: present
 - template: sls_app.yml
 - template_parameters:
 - Environment: “{{ env }}”
 - ...
- cloudformation_facts:
 - stack_name: “{{ env }}-sls-app”

```
- debug:  
  var: cloudformation["{{ env }}-sls-app"].outputs
```

```
[{  
  "description": "Current function version",  
  "output_key": "FuncQualifiedArn",  
  "output_value": "arn:aws:lambda:...:function:MyFunc:3"  
}]
```

Coming Soon

Ansible 2.5 and beyond

AWS Serverless Application Model

- Extensions for CloudFormation server-side transforms
- Rolls packaging/uploading step into stack actions
- Creates a shorthand language for Lambda and API Gateway

Transform: AWS::Serverless-2016-10-31

Resources:

CreateThumbnail:

Type: AWS::Serverless::Function

Properties:

Handler: *handler*

Runtime: *runtime*

Timeout: 60

Policies: AWSLambdaExecute

Events:

CreateThumbnailEvent:

Type: S3

Properties:

Bucket: !Ref SrcBucket

Events: s3:ObjectCreated:*

Sample Play

- cloudformation:
 - stack_name: sam-test-app
 - state: present
 - project_path: “{{ somewhere }}”
 - template: sam_template.yml
- cloudformation_facts:
 - stack_name: sam-test-app

Thank You!

Questions*?

Tweet @ryan_sb or find me at the social tonight

ANSIBLE

-

Serverless

